

IN THE UNITED STATES DISTRICT COURT
FOR THE NORTHERN DISTRICT OF GEORGIA
GAINESVILLE DIVISION

CHRISTOPHER CARRELHA, et al.,

Plaintiffs,

v.

MEB LOAN TRUST VIII, et al.,

Defendants.

CIVIL ACTION FILE

NO. 2:26-cv-00110-RWS-AWH

Judge Richard W. Story

Magistrate Judge Anna W. Howard

FINAL QUALITY GATE — TWO INDEPENDENT AUDITS

AUDIT A — SOURCE, CITATION, FACTUAL, ACCOUNTING AND PROCEDURAL ACCURACY

Run by `audit_accuracy.py`, which is deliberately independent of the validators. The validators ask whether each register is internally well formed. This asks whether the **prose and the registers still agree**, which is where a matter pack drifts once the registers stop changing. It recomputes rather than re-reads.

What it found, and what happened next

Three broken transaction pairings. TXN-010 pointed at TXN-008 and TXN-009 without either pointing back; TXN-023 held the \$855 that entered suspense and never pointed at the row recording it leaving. A relationship recorded from one end only is one a reader traces in the wrong direction and abandons. The data was repaired and `validate_registers.py` now enforces mutuality, which it never had.

A fourth gap the new rule immediately caught. TXN-050 — the May 17, 2025 tender of \$1,899.56 that was never reversed — had no counterpart and no explanation, because there genuinely is no counterpart. That absence is the most significant thing in the ledger, and it was being carried as a blank. The rule was corrected to demand either a named counterpart **or** the word NONE with a stated basis, and the row now records its own absence affirmatively.

A discrepancy that turned out to be the audit's. The first version of this audit recomputed the fee total as \$9,003.42 against the register's \$8,408.74 — same row count, \$594.68 apart. Tracing it

showed the audit was wrong twice: it used its own regex rather than DISP-006's explicitly stated set definition, and it summed **absolute** values, so two fee waivers totalling \$832.52 were added to the charges instead of subtracted from them.

Corrected to sum signed amounts by the ledger's own classification codes, it now **independently confirms** DISP-006:

Classification	Signed total
fee-disbursed	-\$6,968.74
fee-assessed	-\$1,750.90
fee-waived	+\$832.52
Net effect on the account	-\$7,887.12

DISP-006 states \$1,440.00 in SLS-era assessments plus \$6,968.74 in Shellpoint-era disbursements. The disbursement figure reproduces exactly. The assessment figure reproduces exactly once the three late charges ($\$103.39 + \$103.39 + \$104.12 = \310.90) are removed, which is what DISP-006's stated assumptions exclude: $\$1,750.90 - \$310.90 = \$1,440.00$. Both halves of \$8,408.74 are now confirmed by a computation that did not exist when the figure was written.

The net figure of -\$7,887.12 is a different and equally true number, on a different definition. Both belong in the record; neither should be quoted without its definition.

Two work products asserting a court's finding without naming the order. Corrected in WP-001 and WP-002; both now cite Doc. 34 or Doc. 35 inline.

Audit A result

PASS. Zero errors, zero warnings, after four data repairs, two validator corrections and two corrections to the audit itself.

Standing measurements: nineteen sources cited across twenty-one documents, six carrying a reduced access ceiling that every citing document inherits; twenty-three citation records, seven asserting good law, none above medium; seventy-two transactions with thirteen mutual pairings and one movement affirmatively recorded as having no counterpart; four deadlines marked not computable, with no prose anywhere supplying a date for any of them.

AUDIT B — DEFENCE COUNSEL AND A SCEPTICAL MAGISTRATE

Read adversarially, from two seats, against the work product as it now stands.

Seat one — defence counsel

"Your own file says two different things about who owns this loan, and you have read neither instrument." This is the strongest attack available and it lands. The ownership conflict rests on a MERS lookup and a servicer letter. No note, no allonge, no security deed, no assignment instrument,

no trust agreement, no loan schedule has been read. [UNRESOLVED] is the honest label and it is the label the register carries — but a jury instruction cannot be built on it, and neither can a claim.

Holds? Yes. The answer is discovery, not argument. Nothing in the work product overstates this, which is the only reason it is survivable.

"Ames forecloses your assignment theory, and your own materials cite it for the opposite."

Verified and correct. *Ames* holds borrowers lack standing to challenge an assignment of their security deed. The materials in the client's repository cite it as reserving a void-assignment attack. It does not.

Holds? Yes, completely, as to the assignment theory. The corrective action is already taken: the citation is withdrawn in WP-003 on the filer's own initiative, which is the only posture in which this costs credibility once rather than twice.

"You cited *Racette* for a holding it does not contain." Verified. Zero occurrences of "assignment" in the opinion.

Holds? Yes — and it is the more dangerous of the two, because a real case cited for a fictional holding is the fact pattern behind Rule 11 citation inquiries. Withdrawn in WP-003. The recovery is that *Racette* is live authority for a defective-advertisement claim, which is a different and better use of it.

"Your fee complaint is an accounting exercise, not a claim." Partly right. The amounts are the servicer's own and they foot; what is disputed is whether they were authorised, which is a contract and regulation question. The register already classifies DISP-006 as a characterisation dispute rather than an arithmetic one, which is correct and is also the concession.

Holds? Partly. It defeats an overcharge framing and does not defeat a fee-authority framing. GAP-017 — the note and deed provisions authorising each fee type — is what converts one into the other.

"Your limitations problem is unexamined." 15 U.S.C. § 1692k(d) has not been read. The FDCPA count rests on it.

Holds? Yes, and this is the cheapest unfixed exposure in the matter.

"Your client's own strategy documents contain a trust from a different case." Fannie Mae REMIC Trust 2024-91 is a *Caudell* entity. If it reached a filed pleading it is a self-inflicted wound requiring no work from the other side.

Holds? Unknown, and unknowable without the docket. That is itself the finding.

Seat two — a sceptical magistrate

"You were given specific instructions in July. Did you follow them?" The system cannot answer. That is the correct answer for the system to give, and it is also the reason nothing here is ready to file. Against *Vibe Micro* at 1309, this is the question that decides the matter.

"You are asking me to infer wrongdoing from a balance that did not move." The work product does not do this. The unreversed tender is labelled [INFERENCE] and paired with the express

statement that it is consistent with non-application but does not establish it. That discipline is what keeps the strongest fact usable.

"Why are twenty-two attorney decisions unanswered?" Because no human has worked the file. Seventeen of them are marked blocking. A matter in which seventeen blocking questions are open is not a matter ready for a dispositive filing.

"Your deadline register says the repleading deadline passed a month ago and you do not know what happened." Correct, and stated on the face of every document. No date has been invented to fill the hole, which is the one thing that would make it worse.

What Audit B could not test

Whether the substantive Georgia elements are correctly stated, because they remain unresearched (OLQ-7). Whether *Haynes* forecloses the § 44-14-162(b) theory, because it has not been read. Whether any of this survives contact with the actual motion to dismiss, because that motion has not been located.

Audit B result

The work product does not overstate the record anywhere the audit could reach. Its weakest points are weak because the evidence is missing, not because the analysis is wrong — with one exception, now cured: two citations that came into the matter from the client's own repository and would have failed on contact with an opponent.

DEFECTS CORRECTED DURING THIS GATE

#	Defect	Correction
1	Three non-mutual transaction pairings	Data repaired; mutuality now enforced
2	TXN-050's missing counterpart carried as a blank	Recorded as NONE with a stated basis
3	Fee total recomputed wrongly by this audit	Audit corrected to signed sums by classification; register independently confirmed
4	Two court-found assertions with no order named	Both now cite Doc. 34 or Doc. 35
5	<i>Ames</i> cited for a reservation it does not contain	Withdrawn; CON-012 and CV-018 opened
6	<i>Racette</i> cited for a holding it does not contain	Withdrawn; CON-013 and CV-019 opened; correct use identified
7	A trust from a different case in the strategy materials	CON-009 opened; excluded expressly in WP-003
8	Contradictory case number and firm history in the repository	CON-010 and CON-011 opened

SECOND GATE — THE PHASE 5 BUILD

Five further defects, all found by running the new work rather than reading it.

#	Defect	Correction
9	This system produced a non-compliant filing. WP-003 went out in 13-point Times New Roman. N.D. Ga. LR 5.1 permits that typeface only at 14 point or larger. The format profile carried a note saying "read the current rule and set <code>font_size</code> from it" — and nobody did.	The rule was read and recorded (CV-031). <code>produce_document.py</code> now refuses to render a filed document that violates a court type rule it has read, and refuses a profile name it has no rule for. A note is not a check.
10	The docket-ingest tool silently discarded its own output. It wrote row keys invented from memory rather than read from the register header, and the CSV writer dropped every unknown key without a word. It reported "6 rows added" and wrote six blank rows.	Column names corrected, and the writer now refuses to write at all when a populated column is not in the header. A writer that discards data silently turns a typo into undetectable loss.
11	A notice of removal classified as a complaint, because the classifier read the attachment list — "(Attachments: Exhibit A State Court Complaint)".	Classification now reads the opening of the entry, where the clerk names the filing, and falls back to the whole text only if nothing matches.
12	The ingest wrote <code>deadline_type</code> and <code>date_status</code> values that do not exist in the register's enums, so every row it opened failed validation. An ingest that produces rows a human must redo saves nothing.	Corrected to the register's own vocabulary. Every ingested row is now <code>needs-verification</code> with <code>ambiguity_flag: yes</code> , and a period-setting entry becomes a <code>conditional</code> row reading <code>NOT - COMPUTABLE</code> with its trigger named.
13	QC reported "paragraph 3" and "¶ 116" in a motion as dangling internal references. They point at the <i>pleading</i> .	Only Section and Part references are resolved against the document's own headings; the rest are counted and reported as outward.

And the finding that changed the risk picture

Checking the operative pleading directly — rather than reasoning from the strategy documents that discuss it — established that **the two citation defects never reached it**. Zero occurrences of *Racette*, Fannie Mae, REMIC Trust 2024-91, or any void-assignment attack. Paragraph 3 was drafted around *Ames*.

That is the opposite of what the first gate assumed, and it was one command away the entire time. Recorded here because the lesson generalises: a document that *discusses* a filing is not evidence of what the filing says.

Two defects that **are** in the pleading were found in the same pass, and they are the ones that matter — a cross-reference in paragraph 3 pointing at the wrong counts, and one count incorporating "the substantive facts of Counts VI, VIII, XII, and XIII", which contradicts paragraph 2's own representation to the Court and is the second *Weiland* category that Doc. 35 ordered cured.

VALIDATIONS RE-RUN AFTER CORRECTION

All registers; cross-references (422 identifiers defined, 614 references checked, 22 decisions raised and 22 logged); the matter-pack structure with hash-check confirming every original source byte-identical; the access map and credential scan; citation verification with cross-check against the work

product; the approval gate; every handoff document; all eleven skills in both source and installed form; and document QC on every produced document.

Prepared for attorney review,

This 21st day of September, 2026.

Litigation OS — Final Quality Gate

Matter CARRELHA-2026

Not for filing. Not served. Not sent.