

Phase 1

Build, Validation & Forward Test

Four reusable specialist skills, a matter-pack system, an agent handoff standard, and six deterministic validators — built, installed, and forward-tested against the live Carrelha v. MEB Loan Trust VIII record.

4	6	5	5
SKILLS INSTALLED	VALIDATORS	FORWARD TESTS	DEFECTS FIXED

TEST MATTER	COURT
Carrelha v. MEB Loan Trust VIII, et al.	N.D. Ga., Gainesville Div. — 2:26-cv-00110-RWS-AWH
INSTALLED AT	SOURCE OF TRUTH
~/ .claude/skills/	litigation-os/ (repo)
PREPARED	STATUS
21 September 2026	PHASE 1 COMPLETE

Attorney work product — prepared at the direction of counsel. Not legal advice. Nothing has been filed, served, or transmitted. No original case document was modified, renamed, moved, or deleted; source integrity was verified by SHA-256 before and after every run.

1 What was built

The design rule was a system organised by *legal function*, not by lawsuit. The four specialists hold no facts about any case. Each lawsuit is a **matter pack** — a fixed folder of facts, parties, sources, docket, deadlines, issues and evidence — and the specialists operate on whichever pack they are pointed at. Shellpoint was the first test matter; no Shellpoint fact appears in any skill.

The four installed skills

SKILL	FUNCTION	TRIGGERS ON
litigation-matter-orchestrator	Managing attorney. Scopes, routes, stress-tests, reconciles, and signs one consolidated report.	“work plan for this case”, “what should we do next”, “analyze the complaint and the record”, “pull this together”, “case assessment”, “matter status” — and any request needing more than one of research, evidence, or docket work.
legal-research-paralegal	Legal authority. Jurisdiction first, then the hierarchy; verifies eight fields separately; hunts adverse authority.	“what is the law on...”, “does this claim hold up”, “what standard applies”, “is this case still good law”, “check these citations”, “what authority supports/defeats...”.
evidence-chronology-paralegal	Factual record. Inventory, chronology, entity map, proposition/evidence matrix, contradictions, gaps, witness index.	“review these documents”, “build a timeline”, “find the contradictions”, “what does the evidence actually show”, “what’s missing”, “who are the parties”.
docket-deadline-paralegal	Procedural control. Docket register, deadline register, posture, unresolved questions, critical-date warnings.	“what’s our deadline”, “when is this due”, “where does the case stand”, “is this timely”, “what’s been filed”, “appeal period”, “scheduling order”.

THE RULE THAT SHAPES EVERY SKILL

Every substantive statement carries one epistemic label — [VERIFIED] , [ALLEGED] , [COURT-FOUND] , [INFERENCE] , [LEGAL-CONCLUSION] , [UNRESOLVED] — so the attorney never has to guess whether something is proved or merely pleaded. The commonest error the skills are drilled against: a court order reciting “*Plaintiffs allege X*” supports [ALLEGED] , never [COURT-FOUND] .

Files created

PATH	ITEMS	CONTENTS
skills/<each of 4>/	SKILL.md + refs	Operational SKILL.md (190–220 lines each); references carry the long protocols
standards/handoff-standard.md	1	Assignment/result contract, confidence scale, epistemic labels, standing prohibitions
.../references/	9	reconciliation-protocol · work-product-report-template · worked-example · authority-hierarchy · verification-protocol · extraction-protocol · output-schemas · output-formats · deadline-computation-protocol

matter-pack-template/	20 files, 12 dirs	All 18 required components; <code>FIELD-DEFINITIONS.md</code> defines every column so another agent cannot misread one
tools/	6 validators	<code>validate_matter_pack</code> · <code>validate_handoff</code> · <code>validate_registers</code> · <code>validate_crossrefs</code> · <code>validate_skill</code> · <code>common</code>
install.sh	1	Syncs validators + template into each skill, validates, installs to <code>~/.claude/skills/</code>
matters/carrelha-v-meb-2026/	full pack	The populated test matter: 13 sources, 41 events, 13 propositions, 7 contradictions, 12 gaps, 6 deadlines, 18 docket rows, 13 issues, 7 research rows
test-log/defects.md	1	Five corrected defects and four documented-but-uncorrected limits

What the validators refuse to do

DELIBERATE OMISSION

No script computes a legal deadline. Doing so requires the governing rule, the triggering event and its date, the jurisdiction, the service method, and the court’s calendar including local closures. A tool that produced dates from partial inputs would manufacture false confidence — the exact failure the system exists to prevent. Instead the validator enforces the discipline that makes a *human-computed* deadline auditable: a `calculated` row must show its rule, trigger, trigger date, method and calendar basis, or it is rejected; an `estimated` or ambiguous row must read `NOT-COMPUTABLE` and can never be marked confirmed.

2 Forward testing on real Shellpoint materials

All testing used the actual supplied record, reached through the client’s Google Drive and the CourtListener RECAP API. **Originals were read in place and never downloaded, modified, renamed, moved, or deleted.** What is stored locally are verbatim extracts, each registered with its own source ID and a pointer back to its parent, and each SHA-256 fingerprinted.

Materials actually tested

ID	DOCUMENT	READ	WHAT IT CARRIED
SRC -001	First Amended Verified Complaint (draft)	In full	9 counts; the allegations tested against the record
SRC -002	Doc. 34 — Order, 20 Jul 2026 (Story, J.), 26 pp.	In full	Remand denied; PI denied as moot; 60-day foreclosure-notice obligation
SRC -003	Doc. 35 — Order, 31 Jul 2026 (Howard, M.J.), 9 pp.	In full	Shotgun-pleading holding; amend by 21 Aug 2026 ; MTDs moot with leave
SRC -004	Plaintiffs’ First Requests for Production, 19 Aug 2026	In part	Reqs. 1–18; refers to the Rule 26(f) conference in future tense
SRC -005	Shellpoint letter re CFPB 260302-29444848, 2 Mar 2026	In full	Party admissions; designated NOE address; FDCPA disclosure
SRC -006	Shellpoint Loan History Statement, printed 10 Mar 2026, 5 pp.	In full, line by line	The complete transaction ledger, 08/2022 – 02/2026
SRC -007	Annotated ledger spreadsheet with “RED FLAGS” column	In part	Third-party annotations — <i>not</i> a servicer record
SRC -008	CourtListener RECAP docket record, docket id 73194013	In full	Case metadata only — zero entries, zero documents

Test results

TEST	RESULT	WHAT IT PRODUCED
1. Orchestrator bounded work plan	PASS	Reframed five activities as four decisions and <i>reordered</i> them — “is the case still alive?” was not in the ask and became priority one. Three parallel Wave-1 assignments issued, all schema-valid. Routing split the mixed question (“does the amendment cure the defect?”) into a pleading-text half and a legal-standard half rather than sending it whole to one specialist.
2. Research pleading sufficiency	PASS	Opened and read <i>Weiland</i> , 792 F.3d 1313, and <i>Vibe Micro</i> , 878 F.3d 1291, with star pagination confirmed. Found the dispositive holding, the boundary case, adverse authority, and a downside limit nobody asked for. Correctly reported three authorities as unopened or unlocatable, and no conclusion depends on any of them.
3. Evidence manifest + chronology	PASS	13-source manifest, 41-event source-linked chronology, 13 propositions, 7 tested contradictions, 12 gaps, 13 entities. Produced two findings that reframe the matter (below) and corrected an existing error in the client’s own work product.

4. Docket posture + deadlines	PASS	18 docket rows, each with an honest <code>verification_status</code> ; 6 deadlines. Declined to compute four dates for want of inputs and said which input was missing in each case. Flagged the order/rule trigger conflict (“14 days after it is <i>filed</i> ” vs. Rule 15(a)(3)’s <i>service</i>) without silently resolving it.
5. Integrated consolidated report	PASS	232-line attorney work product separating established / alleged / inferred / unverified / evidence needed / viable / weak-or-foreclosed / next steps. Five reconciliations performed across four levels; one of them disclosed the orchestrator’s <i>own</i> planning omission rather than quietly filling it.

Two findings that reframed the matter

FINDING 1 — THE PLEADING’S CENTRAL ALLEGATION IS CONTRADICTED BY THE DEFENDANTS’ OWN RECORD

The First Amended Complaint makes an uncredited **\$5,200.00** wire the factual core of the RESPA, FDCPA and breach-of-contract counts. The servicer’s complete five-page ledger contains **no \$5,200.00 entry in any column, any period, or any servicer era** — and that same ledger records credits as small as \$71.78 and as large as \$19,588.88. This is materially worse than an unproven allegation: the document a factfinder consults first affirmatively omits it.

FINDING 2 — A BETTER PAYMENT FACT, DOCUMENTED, SITTING UNUSED

The same ledger shows four web payments of **\$1,899.56** tendered May–June 2025. Three were reversed with NSF events. **The fourth, 17 May 2025, was never reversed** — and appears as the “Ending Unapplied Balance \$1,899.56” on the 10 March 2026 print. Money received, held in suspense for roughly ten months, never applied, while foreclosure fees accrued against a principal balance frozen since December 2023. Weaker in dollars, far stronger in proof, and drawn from the defendants’ own record rather than from the pleading.

Other substantive results

- **The amendment method appears sound.** *Weiland* at 1324 holds that realleging a block of factual paragraphs, without incorporating preceding *counts*, “looks, at first glance, like the most common type of shotgun pleading. **But it is not.**” That is precisely the draft’s structure — the specific defect Doc. 35 identified appears cured.
- **A curable rendering defect.** Each Count realleges “Paragraphs 1 through 46,” but the numbering restarts in each section, so no paragraph is uniquely 46. Counting continuously the sections run 2+4+7+6+5+7+3+8+4 = **exactly 46** — intent unambiguous, rendering wrong, fixable in an afternoon.
- **A limit on the dismissal risk.** *Vibe Micro*’s one-chance-then-prejudice rule is stated for litigants *represented by counsel*, and footnote 6 expressly “decide[s] and intimate[s] nothing about a party proceeding pro se.” Doc. 35’s own authority for the pro se rule is an unpublished disposition this session could not locate.
- **A downside bound found while hunting adverse authority.** *Vibe Micro* at 1308–09: state-law claims dismissed on non-merits Rule 8 grounds should be dismissed *without* prejudice as to refiling in state court.
- **A correction to existing client work product.** An annotated spreadsheet flags the 5 Aug 2022 row as a “NEGATIVE principal payment” red flag. Read against the original ledger it is the *opening-balance booking*. Advancing it would have been a serious credibility error.

- **Five blocking unknowns, and where they live.** Three are most likely in the client's own files; two are public records obtainable today. Only a minority genuinely require discovery.

3 Defects found by testing, and corrections

Testing is only worth running if it changes the thing being tested. Five defects were found and corrected; four limits were found and documented rather than papered over.

ID	DEFECT	STATUS	CORRECTION
D-01	The handoff validator rejected a correctly-written prohibition because it matched on vocabulary, not on concept.	FIXED	A validator that fails correct input teaches authors to paste boilerplate, converting a safeguard into a ritual. Broadened the matcher, rewrote the error to name the three standing prohibitions, and stated in the standard that they are checked by concept and may be worded to fit the task. The assignment itself was <i>also</i> genuinely incomplete and was reworded.
D-02	Nothing checked internal identifier cross-references. Only SRC- ids were validated.	FIXED	A dangling CLM-014 reads as a mapped element and is not one. Added <code>validate_crossrefs.py</code> , resolving ten identifier families, rejecting malformed ids, ids of the wrong family for their column, and duplicate definitions. Now checks 247 references against 123 definitions in the test matter.
D-03	The confidence rule made high unreachable whenever no citator exists — collapsing a three-level scale to two.	FIXED	The validator was right and the rows were wrong; four research rows were downgraded to medium. But the rule was <i>also</i> hiding a distinction worth keeping, so the protocol now names citator-limited medium : everything checkable against the opinion was checked, and what is owed is a Shepard's/KeyCite pass, not more reading. An attorney can now tell those two situations apart at a glance.
D-04	Register confidence caps did not propagate into narrative prose.	FIXED	Caught by an audit script, not by reading. A register row correctly capped at medium fed a paragraph written as certainty. The attorney reads the paragraph. Added the inheritance rule to the standard and a mandatory readback pass to the evidence skill; restated the finding in the report rather than editing it away.
D-05	The orchestrator planned Wave 1 and never recorded the dependent Wave 2, losing a task.	FIXED	The orchestrator's own error. A dependent task held only in the planner's head goes missing exactly when results arrive and attention shifts to reconciling them. Step 3 now requires dependent tasks to be written into the issue register at plan time with their dependency named; Step 7 requires every planned task to be accounted for. Four Wave-2 tasks are now recorded as blocked in the test matter.

Limits documented, not corrected

ID	LIMIT	WHY IT STANDS
N-01	No validator can confirm a quotation is accurate.	The validator has no access to the source. Mitigated by requiring a pinpoint and a named verification method; not eliminated. A human still has to check quotations.
N-02	Epistemic labels are enforced for membership, not correctness.	A script can reject a misspelled label; it cannot tell that a passage marked [COURT-FOUND] is really a court reciting a party's allegation. Addressed by instruction in four places, because no validator can reach it.

N-03	The system cannot see PACER.	The most important fact in the test matter was unobtainable. The system reported it as a blocking human decision rather than inferring an answer — designed behaviour, standing limit.
N-04	Empty registers pass validation.	Deliberate. A matter legitimately has no contradictions until someone looks. An emptiness check would force placeholder rows, which are worse than none. The orchestrator's completion standard, not the validator, requires population.

4 Completion, weaknesses, and Phase 2

Phase 1 completion standard

REQUIREMENT	STATUS	EVIDENCE
All four skills installed	MET	~/.claude/skills/ – all four live
Each skill passes structural validation	MET	4/4 pass, source and installed copies
Matter-pack templates exist	MET	All 18 components + field definitions
Handoff formats exist	MET	Assignment + result, schema-enforced
All four specialists forward-tested	MET	Tests 1–4 on real materials
Integrated workflow tested	MET	Test 5 — 5 reconciliations
Defects corrected or documented	MET	5 corrected, 4 documented
Original case materials untouched	MET	SHA-256 verified; originals never downloaded

Remaining weaknesses

1. **No PACER, and it mattered.** The single most important fact about the test matter — whether the amended complaint was filed by 21 Aug 2026 — could not be established. Until a human answers it, parts of the case assessment are explicitly conditional.
2. **No citator.** Good-law status is unconfirmed on every authority. Every research row is citator-limited medium by design.
3. **Quotation accuracy and label correctness are human-checked.** N-01 and N-02 are the two places where the system’s guarantees stop.
4. **Single-matter testing.** Everything was exercised against one matter in one posture. Matter-agnosticism is a design property enforced by review, not yet demonstrated across a second matter — the cheapest next proof would be a second pack in a different court.
5. **Three authorities remain unread** in the test matter (*Magluta*, *Barmapov*, *Lee v. Alachua County*), all correctly flagged, none load-bearing.
6. **The evidence specialist’s OCR discipline was under-exercised.** Every source was native text. The ambiguous-column handling worked, but true image-only OCR was never tested.

RECOMMENDED PHASE 2 BUILD — IN THIS ORDER

- 1 • pleading-amendment-analyst.** Element-by-element testing of a pleading against a repleading order and the four *Weiland* categories; a pre-filing defect checklist. Highest value because the test matter is *in* the repleading phase and the caution is dismissal.
- 2 • loan-accounting-analyst.** Ledger reconciliation, suspense tracing, payment application, fee authorization, interest recomputation. The \$1,899.56 finding shows the ledger is the richest unexploited source; this specialist turns row-reading into provable arithmetic.
- 3 • adversarial-redteam-reviewer.** Takes a finished work product and argues the other side: attacks each proposition, finds the adverse authority the researcher missed, and rates what survives. Placed third deliberately — it needs finished product to attack, and D-04 showed that self-audit catches what self-review does not.
- 4 • discovery-planner.** Turns the missing-evidence register into requests, subpoenas and public-record pulls, sequenced against deadlines and the Rule 26(d)/26(f) gate. Depends on a populated gap register, which Phase 1 now produces.
- 5 • securitization-ownership-analyst.** Chain of title, note indorsements, assignment validity, trust cut-off dates, standing to enforce the power of sale. Placed last because it is blocked on documents nobody has yet obtained — build it when the recorded instruments are in hand, not before.

Phase 2 has not been started. This report closes Phase 1.