

Phase 2

Five Specialists, Nine Integrated

Phase 1 audited and repaired, five new specialists built and installed, the orchestrator rewired for nine, and the whole system forward-tested against the live Carrelha v. MEB Loan Trust VIII record.

9	26	516	9
SKILLS INSTALLED	REGISTERS POPULATED	CROSS-REFS RESOLVING	DEFECTS FIXED

TEST MATTER

Carrelha v. MEB Loan Trust VIII, et al.

SOURCES

19 registered · 7 read in full · 3 in part

PREPARED

21 September 2026

COURT

N.D. Ga., Gainesville Div. — 2:26-cv-00110-RWS-AWH

WORK PRODUCTS

10 integrated deliverables

STATUS

PHASE 2 COMPLETE

Attorney work product. Not legal advice. Nothing filed, served, or transmitted. No original case document modified, renamed, moved or deleted — SHA-256 verified before and after every run. Phase 3 not begun.

1 Phase 1 audited before anything was built on it

Validators passing is necessary, not sufficient. The Phase 1 deliverables were inspected directly — drift across copies, row counts, word counts, stale paths, hardcoded lists — rather than trusted from the Phase 1 summary. Five defects were found.

ID	DEFECT	STATUS	REPAIR
P1-R01	Two bundled files still gave commands as <code>litigation-os/tools/...</code> . A skill copied to a machine with no repo would hand the user four commands that cannot run.	FIXED	Repointed, and <code>install.sh</code> now greps for repo paths in bundled resources and fails the build.
P1-R02	The handoff validator hardcoded three specialist names. Every Phase 2 assignment would have been rejected, and the natural workaround — editing six copies — is exactly the drift <code>install.sh</code> exists to prevent.	FIXED	Replaced with <code>standards/specialists.yaml</code> , a single registry the validator, orchestrator and installer all read. Adding a tenth specialist is now a one-file change.
P1-R03	Eleven human decisions were raised across the Phase 1 results and none was logged . The decision log is the gate on outward action; empty, there is no record of what anyone authorised.	FIXED	All eleven logged with <code>decided_by</code> blank; a new check reconciles every <code>HD-</code> raised against the log and fails on any missing. It caught 11 more during Phase 2.
P1-R04	Byte-compiled Python committed.	FIXED	Removed; <code>.gitignore</code> added.
P1-R05	A template artefact is copied into every live matter.	ACCEPTED	Deliberate. The validator already skips it, and special-casing the copy adds a moving part to prevent a non-problem.

THE ONE THAT MATTERED

P1-R03 is the substantive one. Phase 1's orchestrator instructions *said* to log human decisions; nothing checked that they were. Eleven requests — five of them blocking — evaporated between the specialists and the attorney. The repair was not just to log them but to make the gap **mechanically detectable**, because a rule stated in prose and enforced by nothing is a rule that will be broken exactly when it matters.

2 The five new specialists

SKILL	WHAT IT DOES	THE RULE THAT MAKES IT SAFE
pleading-amendment-analyst	Element-by-element, defendant-by-defendant claim mapping; a 14-item defect scan; Rule 15 analysis; claim-survival matrix, amendment memo, correction list, pleading-support table.	Never proposes an allegation without naming its basis — a source and pinpoint, a named witness, or an explicit [BASIS-REQUIRED] flag. A correction list reading "allege that Defendant knew" invites a sanctionable filing.
loan-accounting-analyst	Transaction-level reconciliation; classification into twelve categories with pairing; disputed-amount schedule; unexplained-variance report.	Preserves the original and never manufactures ledger logic. A label stays verbatim even when it contradicts the entry's own effect. An unexplained entry is not proof of wrongdoing.
discovery-planning-analyst	Gaps → targeted, defendant-specific requests; roadmap, custodian map, request-to-issue matrix, deposition plan, objection analysis.	Drafts only — never serves. Every request traces to a gap <i>and</i> an element, or it is cut. The gate is checked before planning, not after.
securitization-ownership-analyst	Chain tracing across originator, note, allonge, assignments, servicer, trust; entity map, transfer chronology with eight date types, authority matrix.	Separates material defects from theories courts reject , and keeps ownership, enforcement, servicing and reporting strictly apart. The rejected-theory list is a required deliverable, not an appendix.
adversarial-redteam-analyst	Five independent seats: defense counsel, skeptical magistrate, district judge, Rule 11 reviewer, appellate preservation. Attack surface, ranking, credibility risks, curability, corrective actions.	Argues to win, for them. Concedes what is true — an honest red team is only valuable if its "this one holds up" is credible.

Integration: nine specialists without flooding

Nine specialists can generate more paper than any attorney will read, and can duplicate each other's work in ways that manufacture spurious conflicts. The orchestrator gained three things:

- **A tier system.** Tier 1 reads raw sources and starts immediately. Tier 2 consumes Tier 1. Tier 3 consumes Tier 2 — discovery planning needs proven gaps, a red team needs finished work to attack. Deviating is allowed; doing it silently is not.
- **A no-duplication rule.** Never two assignments whose `required_output` overlaps. If two specialists need the same artefact, one produces it and the other is given it as a source. Overlapping *reading* is fine; overlapping *output* is waste.
- **A four-axis ranking step.** Urgency, legal significance, evidentiary strength, curability — with explicit heuristics: urgency beats significance; cheap-and-curable beats expensive-and-severe; *contradicted* beats merely unsupported.

And a flooding rule with teeth: before issuing an assignment, answer *which decision does this change?* If the answer is "it would be good to know," do not issue it. Say in the plan which specialists you are **not** using, and why.

3 Shellpoint forward test

Nineteen sources registered. Originals were read in place in the client's Drive and **never downloaded, modified, renamed or deleted**; what is stored locally are verbatim extracts, each with its own source ID pointing back at its parent, each SHA-256 fingerprinted.

Sources actually reviewed

ID	DOCUMENT	READ	WHAT IT CARRIED
SRC - 020	First Amended Complaint — 16 counts, 2026-08-19	In full	A different and later pleading than Phase 1 analysed
SRC - 006/013	Shellpoint Loan History Statement, printed 2026-03-10, 5 pp.	Every row	All 72 transactions, 08/2022–02/2026
SRC - 024	File named "5200 on Ledger", printed 2026-01-16	In full	A second print of the same statement — and no \$5,200 entry
SRC - 002/003	Doc. 34 (26 pp.) and Doc. 35 (9 pp.)	In full	The two orders; the 60-day notice obligation; the 8/21 deadline
SRC - 005	Shellpoint CFPB response, 2026-03-02	In full	Debt-collector admission; the designated NOE address
SRC - 022	DBRS press release — FREED Mortgage Trust 2022-HE1	In part	Cut-off 2022-09-30; 3,292 Lendage HELOCs; sponsor
SRC - 001	First Amended Verified Complaint — 9 counts	In full	The Phase 1 draft, now superseded
SRC - 004	Plaintiffs' First Requests for Production, 2026-08-19	In part	Rule 26(f) referred to in future tense
SRC - 007	Annotated ledger spreadsheet, "RED FLAGS" column	In part	Litigant work product — not a servicer record
SRC - 008	CourtListener RECAP docket record	In full	Metadata only — zero entries, zero documents

NOT REVIEWED — AND THE ANALYSIS SAYS SO EVERYWHERE IT MATTERS

No **primary ownership document** was read: no note, allonge, endorsement, assignment instrument, security deed, trust agreement, or mortgage loan schedule. The ownership result carries **confidence: low** for exactly this reason and says so in its first line. No credit report, no ACDV record, no NOE or RFI, no loan modification, no employment document, no bank record, and no PACER docket. **Counts X–XVI of sixteen were read at heading level only.**

What the test found

FINDING 1 — PHASE 1 ANALYSED THE WRONG PLEADING

A later 16-count First Amended Complaint exists, dated 2026-08-19. It is materially better on every axis the two orders put in issue: **no count incorporates a preceding count**; numbering is continuous; **Lendage is out of the caption**; the **\$5,200 allegation is gone**; and ¶ 3 is drafted expressly around the Georgia standing rule. Phase 1 had flagged this risk as HD-2. The risk materialised. **Phase 1's pleading analysis is superseded.**

FINDING 2 — THE \$5,200 WIRE, REFINED RATHER THAN REVERSED

Phase 1 reported no \$5,200 entry exists. Still true. With the December 19, 2023 date now supplied, the 12/21/23 cluster becomes a candidate — and equals **\$4,355.87**, **\$5,085.00**, or **\$4,980.88** depending on which entries are counted. **None is \$5,200.00**. All three computations are shown rather than one chosen silently. The honest answer is neither "absent" nor "applied" — it is unexplained, and the later pleading drops the allegation entirely.

FINDING 3 — THE PLEADING UNDERSTATES ITS OWN BEST EVIDENCE

¶ 29 pleads that *each* of four 2025 trial payments was reversed. The ledger Plaintiffs rely on as their own exhibit shows **three of four**: the **2025-05-17 tender was never reversed** and equals the ledger's own ending unapplied balance of \$1,899.56, ten months later. An exhibit contradicting an allegation is the cheapest point a defendant can make — and here the correction makes the allegation *stronger*. Ten minutes to fix.

FINDING 4 — THE DEFENDANTS' OWN RECORDS DISAGREE ABOUT WHO OWNS THE LOAN

The MERS registry names **FREED Mortgage Trust 2022-HE1** as investor as of 2026-03-24. Shellpoint's own letter names **MEB Loan Trust VIII** as owner. The recorded assignment recites MEB as "a[n] **CORPORATION organized under the laws of NEW YORK**" while the letter gives a Delaware trust address. **Its value is not what it first appears**: pleaded as a chain-of-title attack it meets the standing bar and loses; pleaded as a records-reliability point it raises no standing question at all, because it does not challenge the assignment.

Other results

- **¶ 31 sues the wrong defendant.** It attributes an \$855 rejection to Shellpoint; the ledger marks both rows **PriorServicer**, ten months before Shellpoint's boarding — which ¶ 27 itself pleads. Provable from our own exhibit. Five minutes to fix.
- **The Doc. 35 defect is cured.** 6 of 14 defect tests return clean, including the type-one shotgun test and the ambiguous cross-reference that afflicted the earlier draft.
- **\$8,408.74** in 27 default and foreclosure charges against a principal frozen at \$140,288.02 since December 2023 — the pleading claims "at least fifteen... exceeding \$3,500" and *understates* the record.
- **The reinstatement reconciles to within \$70.00.** \$19,588.88 in on 2023-09-13, \$19,518.88 distributed the next day. Small, unexplained, and reported as unexplained — not as wrongdoing.
- **An error in the client's own work product corrected.** An annotated spreadsheet flags the 2022-08-05 row as a "NEGATIVE principal payment" red flag. It is the opening-balance booking. Advancing it would be refuted in one sentence.
- **Five of eight blocking gaps need no discovery at all** — three are in the client's files or their bank's, two are public records.

Ten integrated work products

#	DELIVERABLE	WHERE
1	Pleading and amendment survival matrix	claim-survival-matrix.csv (13) + pleading-defects.csv (14)
2	Transaction-level accounting reconciliation	transaction-reconciliation.csv (72) + disputed-amounts.csv (7)
3	Ownership, funding, transfer and servicing chronology	entity-role-map.csv (14) + transfer-chronology.csv (6)
4	Proposition-to-evidence matrix	proposition-evidence.csv (13) + pleading-support-table.csv (13)
5	Contradiction and admission register	contradiction-register.csv (8)
6	Controlling and adverse authority table	research-table (7) + authority-matrix.csv (6)
7	Discovery roadmap	request-to-issue.csv (9) + custodian-map.csv (7)
8	Defense and judicial red-team report	attack-surface.csv (10) + corrective-actions.csv (10)
9	Missing-evidence priority list	missing-evidence.csv (21)
10	Consolidated attorney work-product memorandum	MEMO-CARRELHA-2026-002-integrated.md (247 lines)

4 Validation, defects, and Phase 3

Validation results

CHECK	RESULT	DETAIL
Skill structure — all nine, source and installed copies	9/9	Registry/install/disk agreement enforced at build time
Blank matter-pack template	PASS	30 required files, 13 directories
Test matter pack + source integrity	PASS	SHA-256: all originals byte-identical
Registers — all nine kinds	PASS	26 registers populated
Internal cross-references	PASS	299 identifiers, 516 references, all resolving
Human-decision reconciliation	PASS	22 raised, 22 logged, 14 blocking, <code>decided_by</code> blank on all
Handoff documents	16/16	8 assignments, 8 results
Negative tests — 16 defects injected across 7 registers	16/16	21 errors raised; the validators are not vacuous

Defects found by Phase 2 testing

ID	DEFECT	STATUS	CORRECTION
P2-D01	A reversal was declared with no pairing row — the row it reversed had never been transcribed.	FIXED	TXN-072 added and paired. The validator caught an incomplete transcription a reader would not have noticed.
P2-D02	Improvement pass 1. An ID-shaped token with an unregistered family validated clean. <code>CLMM-001</code> — a typo for <code>CLM-001</code> — returned PASS.	FIXED	Anything matching <code>^[A-Z]{2,6}-\d{1,4}\$</code> is now an identifier; an unregistered family is an error naming the token and the known families. It immediately caught three real errors in live data.
P2-D03	Improvement pass 2. The pleading-defect register was entirely unchecked: <code>present: maybe</code> , <code>severity: catastrophic</code> , and a present defect with no cure all validated clean.	FIXED	Four controlled-value checks; <code>pleading_analysed</code> required on every row; a cure, a cost and a curability rating required on every present defect; a warning when fewer than 14 tests are recorded.
P2-D04	The corrective-action plan existed only as a table inside a result. Prose is not addressable — nothing could point at CA-03, and no validator could check that a top-ranked attack had an action behind it.	FIXED	<code>corrective-actions.csv</code> added to template, pack, skill assets and resolver, with 10 rows tracing to the attacks and defects they address.

THE LESSON THAT GENERALISES

P2-D04 is the one worth carrying forward: **an output that exists only as prose cannot be validated, cross-referenced, or tracked to completion.** Where a deliverable is a list of things someone must do, it belongs in a register. The same reasoning drove the P1-R03 repair — a rule stated in prose and enforced by nothing is a rule that breaks exactly when it matters.

Unresolved access and evidence gaps

- **No PACER.** Whether the amended complaint was filed by 2026-08-21 — 31 days past, under an order cautioning dismissal — is still unknown, and **two candidate pleadings exist.**
- **No citator.** Good-law status unconfirmed on every authority.
- **Element formulations are UNRESEARCHED** throughout the claim-survival matrix. A wrong formulation would silently corrupt every row resting on it. This is stated in the matrix, the result, and the memorandum.
- **No primary ownership document was read.** Two of them are public records obtainable today for a nominal fee.
- **Six of sixteen counts unanalysed.** The attack surface is incomplete by construction.
- **No credit report, ACDV record, NOE/RFI, loan modification, employment document or bank record** is in the manifest.
- **A [COL?] figure is excluded row-locally, not globally** (P2-N01) — the guard is the required **assumptions** column plus the discipline of saying so.

RECOMMENDED PHASE 3 BUILD — IN THIS ORDER

1 · motion-practice-analyst. Briefing: opposition and reply structure, standard of review, record citation, local formatting and page limits, proposed orders. Highest value because the matter's next event is a renewed motion to dismiss and Phase 2 has already mapped the attack surface it must answer.

2 · authority-verification-service. A dedicated citator pass: Shepardize/KeyCite every authority, resolve every **UNRESEARCHED** element formulation, and fill the good-law column. Second because it unblocks the single most pervasive caveat in Phase 2 — *every* element list in the claim-survival matrix currently rests on the pleading's own recitals.

3 · damages-quantification-analyst. Statutory, actual and special damages computed element by element with its own evidentiary support. Third because the largest figure in the matter rests entirely on an allegation, and because it needs the accounting output Phase 2 now produces.

4 · settlement-valuation-analyst. Expected-value analysis across outcomes, informed by the red team's ranking and the damages model. Fourth because it depends on both.

5 · filing-qc-and-compliance-analyst. Pre-filing check against local rules, standing orders, formatting, exhibits, signatures and certificates of service. Last because it operates on a finished draft — and because Phase 2 found an unfilled signature block in one candidate pleading and could not assess the other.

Phase 3 has not been started. This report closes Phase 2.